

AI Lab: Deep Learning for Computer Vision

WorldQuant University

Usage Guidelines

This file is licensed under [Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International](#).

You **can**:

- ✓ Download this file
- ✓ Post this file in public repositories

You **must always**:

- ✓ Give credit to [WorldQuant University](#) for the creation of this file
- ✓ Provide a [link to the license](#)

You **cannot**:

- ✗ Create derivatives or adaptations of this file
- ✗ Use this file for commercial purposes

Failure to follow these guidelines is a violation of your terms of service and could lead to your expulsion from WorldQuant University and the revocation of your certificate.

Getting Started

Let's import the packages we'll need in this notebook. Most are familiar. We will need version 2 of the `torchvision.transforms` module here. The API is slightly different than that of version 1 that we've used previously, but it's pretty similar.

```
In [1]: import pathlib
import sys

import matplotlib.pyplot as plt
import torch
import torchinfo
import torchvision
import ultralytics
from PIL import Image
from torchvision.transforms import v2
from ultralytics import YOLO
```

```
Creating new Ultralytics Settings v0.0.6 file ✓
View Ultralytics Settings with 'yolo settings' or at '/root/.config/Ultralytics/settings.json'
Update Settings with 'yolo settings key=value', i.e. 'yolo settings runs_dir=path/to/dir'. For help see https://docs.ultralytics.com/quickstart/#ultralytics-settings.
```

In case we want to reproduce this notebook in the future, we'll record the version information.

```
In [2]: print("Platform:", sys.platform)
print("Python version:", sys.version)
print("---")
print("matplotlib version:", plt.matplotlib.__version__)
print("PIL version:", Image.__version__)
print("torch version:", torch.__version__)
print("torchvision version:", torchvision.__version__)
print("ultralytics version:", ultralytics.__version__)
```

```
Platform: linux
Python version: 3.11.0 (main, Nov 15 2022, 20:12:54) [GCC 10.2.1 20210110]
---
matplotlib version: 3.9.2
PIL version: 10.2.0
torch version: 2.2.2+cu121
torchvision version: 0.17.2+cu121
ultralytics version: 8.3.27
```

These are the classes of the Dhaka AI data set we've seen before.

```
In [3]: CLASS_DICT = dict(
    enumerate(
        [
            "ambulance",
            "army vehicle",
            "auto rickshaw",
            "bicycle",
            "bus",
            "car",
            "garbagevan",
            "human hauler",
            "minibus",
            "minivan",
            "motorbike",
            "pickup",
            "policecar",
            "rickshaw",
            "scooter",
            "suv",
            "taxi",
            "three wheelers (CNG)",
            "truck",
            "van",
            "wheelbarrow",
        ]
    )
)
```

```
)  
  
print("CLASS_DICT type, ", type(CLASS_DICT))  
CLASS_DICT  
  
CLASS_DICT type, <class 'dict'>  
  
Out[3]: {0: 'ambulance',  
         1: 'army vehicle',  
         2: 'auto rickshaw',  
         3: 'bicycle',  
         4: 'bus',  
         5: 'car',  
         6: 'garbagevan',  
         7: 'human hauler',  
         8: 'minibus',  
         9: 'minivan',  
        10: 'motorbike',  
        11: 'pickup',  
        12: 'policecar',  
        13: 'rickshaw',  
        14: 'scooter',  
        15: 'suv',  
        16: 'taxi',  
        17: 'three wheelers (CNG)',  
        18: 'truck',  
        19: 'van',  
        20: 'wheelbarrow'}
```

Data Augmentation in Our YOLO Model

In the previous notebook, we passed our training images to the YOLO model and let it do its thing. The obvious assumption to make is that these images would be used as is, but it turns out not to be so. To demonstrate what was happening, we'll load the model back up and poke around inside of it a bit.

Let's start by finding a saved version of the model. This cell should show all of the training runs that have been completed.

```
In [4]: runs_dir = pathlib.Path("runs", "detect")  
list(runs_dir.glob("train*"))
```

```
Out[4]: [PosixPath('runs/detect/train')]
```

If you don't see anything listed here, go back and run the previous notebook all the way through!

Task 3.5.1: Choose a training run, and check that there are model weights saved in the `weights/best.pt` file for that run.

```
In [5]: run_dir = runs_dir / "train"
weights_file = run_dir / "weights" / "best.pt"

print("Weights file exists?", weights_file.exists())
```

Weights file exists? True

Task 3.5.2: Load the model from the weights file.

```
In [6]: model = YOLO(weights_file)

torchinfo.summary(model)
```

```
Out[6]: =====
Layer (type:depth-idx)           Param #
=====
YOLO
|---DetectionModel: 1-1          --
|   |---Sequential: 2-1          --
|   |   |---Conv: 3-1            (464)
|   |   |---Conv: 3-2            (4,672)
|   |   |---C2f: 3-3             (7,360)
|   |   |---Conv: 3-4             (18,560)
|   |   |---C2f: 3-5             (49,664)
|   |   |---Conv: 3-6             (73,984)
|   |   |---C2f: 3-7             (197,632)
|   |   |---Conv: 3-8             (295,424)
|   |   |---C2f: 3-9             (460,288)
|   |   |---SPPF: 3-10            (164,608)
|   |   |---Upsample: 3-11          --
|   |   |---Concat: 3-12          --
|   |   |---C2f: 3-13             (148,224)
|   |   |---Upsample: 3-14          --
|   |   |---Concat: 3-15          --
|   |   |---C2f: 3-16             (37,248)
|   |   |---Conv: 3-17             (36,992)
|   |   |---Concat: 3-18          --
|   |   |---C2f: 3-19             (123,648)
|   |   |---Conv: 3-20             (147,712)
|   |   |---Concat: 3-21          --
|   |   |---C2f: 3-22             (493,056)
|   |   |---Detect: 3-23            (755,407)
=====

Total params: 3,014,943
Trainable params: 0
Non-trainable params: 3,014,943
=====
```

We need to get the model set up to load the data. The easiest way to do that is to train it for an epoch.

When you call `.train()` on a YOLO model, it sets up a data loader, if it doesn't already exist. Unfortunately, there's no easy way to trigger that set-up step without doing an

epoch of training. 😊

Task 3.5.3: Run one epoch of training.

```
In [7]: result = model.train(  
    data=model.overrides["data"],  
    epochs=1,  
    batch=8,  
    workers=1,  
)
```

Ultralytics 8.3.27 🎨 Python-3.11.0 torch-2.2.2+cu121 CUDA:0 (Tesla T4, 14918MiB)
engine/trainer: task=detect, mode=train, model=runs/detect/train/weights/best.pt, data=data.yaml, epochs=1, time=None, patience=100, batch=8, imgsz=640, save=True, save_period=-1, cache=False, device=None, workers=1, project=None, name=train2, exist_ok=False, pretrained=True, optimizer=auto, verbose=True, seed=0, deterministic=True, single_cls=False, rect=False, cos_lr=False, close_mosaic=10, resume=False, amp=True, fraction=1.0, profile=False, freeze=None, multi_scale=False, overlap_mask=True, mask_ratio=4, dropout=0.0, val=True, split=val, save_json=False, save_hybrid=False, conf=None, iou=0.7, max_det=300, half=False, dnn=False, plots=True, source=None, vid_stride=1, stream_buffer=False, visualize=False, augment=False, agnostic_nms=False, classes=None, retina_masks=False, embed=None, show=False, save_frames=False, save_txt=False, save_conf=False, save_crop=False, show_labels=True, show_conf=True, show_boxes=True, line_width=None, format=torchscript, keras=False, optimize=False, int8=False, dynamic=False, simplify=True, opset=None, workspace=4, nms=False, lr0=0.01, lrf=0.01, momentum=0.937, weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=7.5, cls=0.5, df1=1.5, pose=12.0, kobj=1.0, label_smoothing=0.0, nbs=64, hsv_h=0.015, hsv_s=0.7, hsv_v=0.4, degrees=0.0, translate=0.1, scale=0.5, shear=0.0, perspective=0.0, flipud=0.0, fliplr=0.5, bgr=0.0, mosaic=1.0, mixup=0.0, copy_paste=0.0, copy_paste_mode=flip, auto_augment=randaugment, erasing=0.4, crop_fraction=1.0, cfg=None, tracker=botsort.yaml, save_dir=runs/detect/train2

	from	n	params	module	a
arguments					
0		-1 1	464	ultralytics.nn.modules.conv.Conv	
[3, 16, 3, 2]					
1		-1 1	4672	ultralytics.nn.modules.conv.Conv	
[16, 32, 3, 2]					
2		-1 1	7360	ultralytics.nn.modules.block.C2f	
[32, 32, 1, True]					
3		-1 1	18560	ultralytics.nn.modules.conv.Conv	
[32, 64, 3, 2]					
4		-1 2	49664	ultralytics.nn.modules.block.C2f	
[64, 64, 2, True]					
5		-1 1	73984	ultralytics.nn.modules.conv.Conv	
[64, 128, 3, 2]					
6		-1 2	197632	ultralytics.nn.modules.block.C2f	
[128, 128, 2, True]					
7		-1 1	295424	ultralytics.nn.modules.conv.Conv	
[128, 256, 3, 2]					
8		-1 1	460288	ultralytics.nn.modules.block.C2f	
[256, 256, 1, True]					
9		-1 1	164608	ultralytics.nn.modules.block.SPPF	
[256, 256, 5]					
10		-1 1	0	torch.nn.modules.upsampling.Upsample	
[None, 2, 'nearest']					
11		[-1, 6] 1	0	ultralytics.nn.modules.conv.Concat	
[1]					
12		-1 1	148224	ultralytics.nn.modules.block.C2f	
[384, 128, 1]					
13		-1 1	0	torch.nn.modules.upsampling.Upsample	
[None, 2, 'nearest']					
14		[-1, 4] 1	0	ultralytics.nn.modules.conv.Concat	
[1]					
15		-1 1	37248	ultralytics.nn.modules.block.C2f	
[192, 64, 1]					
16		-1 1	36992	ultralytics.nn.modules.conv.Conv	

```
[64, 64, 3, 2]
17           [-1, 12] 1      0 ultralytics.nn.modules.conv.Concat
[1]
18           -1 1 123648 ultralytics.nn.modules.block.C2f
[192, 128, 1]
19           -1 1 147712 ultralytics.nn.modules.conv.Conv
[128, 128, 3, 2]
20           [-1, 9] 1      0 ultralytics.nn.modules.conv.Concat
[1]
21           -1 1 493056 ultralytics.nn.modules.block.C2f
[384, 256, 1]
22           [15, 18, 21] 1 755407 ultralytics.nn.modules.head.Detect
[21, [64, 128, 256]]
Model summary: 225 layers, 3,014,943 parameters, 3,014,927 gradients, 8.2 GFLOPs
```

Transferred 355/355 items from pretrained weights

Freezing layer 'model.22.dfl.conv.weight'

AMP: running Automatic Mixed Precision (AMP) checks...

AMP: checks passed ✓

```
train: Scanning /app/data_yolo/labels/train... 2422 images, 0 backgrounds, 0 corrupt: 100%|██████████| 2422/2422 [00:00<00:00, 2597.89it/s]
```

```
train: New cache created: /app/data_yolo/labels/train.cache
```

```
val: Scanning /app/data_yolo/labels/val... 578 images, 0 backgrounds, 0 corrupt: 100%|██████████| 578/578 [00:00<00:00, 2344.13it/s]
```

```
val: New cache created: /app/data_yolo/labels/val.cache
```

```
/usr/local/lib/python3.11/site-packages/torch/utils/data/dataloader.py:558: UserWarning: This DataLoader will create 2 worker processes in total. Our suggested max number of worker in current system is 1, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.
```

```
warnings.warn(_create_warning_msg(
```

Plotting labels to runs/detect/train2/labels.jpg...

```
optimizer: 'optimizer=auto' found, ignoring 'lr0=0.01' and 'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum' automatically...
```

```
optimizer: AdamW(lr=0.0004, momentum=0.9) with parameter groups 57 weight(decay=0.0), 64 weight(decay=0.0005), 63 bias(decay=0.0)
```

Image sizes 640 train, 640 val

Using 1 dataloader workers

Logging results to runs/detect/train2

Starting training for 1 epochs...

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
1/1	0G	1.168	1.218	1.048	95	640: 100%
	303/303 [02:29<00:00, 2.03it/s]	Class	Images	Instances	Box(P	R

	all	578	4590	0.613	0.363	0.414
0.286						

1 epochs completed in 0.048 hours.

Optimizer stripped from runs/detect/train2/weights/last.pt, 6.2MB

Optimizer stripped from runs/detect/train2/weights/best.pt, 6.2MB

Validating runs/detect/train2/weights/best.pt...

Ultralytics 8.3.27 🚀 Python-3.11.0 torch-2.2.2+cu121 CUDA:0 (Tesla T4, 14918MiB)

Model summary (fused): 168 layers, 3,009,743 parameters, 0 gradients, 8.1 GFLOPs

	Class	Images	Instances	Box(P)	R	mAP50	mAP50
-95): 100% ██████████ 37/37 [00:16<00:00, 2.30it/s]							
0.286	all	578	4590	0.61	0.366	0.414	
0.601	ambulance	17	17	1	0	0.00737	0.0
0.359	army vehicle	10	12	1	0.164	0.427	
0.322	auto rickshaw	37	111	0.557	0.486	0.523	
0.155	bicycle	72	93	0.408	0.28	0.326	
0.515	bus	295	645	0.679	0.69	0.73	
0.517	car	309	1006	0.748	0.677	0.74	
0.145	human hauler	23	25	0.368	0.24	0.201	
0.286	minibus	10	11	0	0	0.0416	0.
0.255	minivan	128	208	0.387	0.442	0.352	
0.345	motorbike	238	449	0.647	0.561	0.601	
0.38	pickup	157	217	0.743	0.429	0.558	
0.166	policecar	6	6	1	0	0.223	
0.457	rickshaw	202	653	0.475	0.749	0.706	
0.049	scooter	8	9	1	0	0.00812	0.
0.271	suv	105	161	0.274	0.514	0.366	
0.435	taxi	8	10	1	0	0.55	
0.483	three wheelers (CNG)	224	510	0.686	0.658	0.697	
0.502	truck	159	289	0.558	0.73	0.694	
0.164	van	86	139	0.302	0.432	0.246	
0.203	wheelbarrow	16	19	0.359	0.263	0.279	

Speed: 0.6ms preprocess, 11.3ms inference, 0.0ms loss, 4.0ms postprocess per image

Results saved to `runs/detect/train2`

The model should now have a `.trainer` attribute, which has a `.train_loader` attribute. This will be a `DataLoader` that loads the training data.

Task 3.5.4: Save this data loader to variable `loader`.

```
In [10]: loader = model.trainer.train_loader
print(type(loader))
<class 'ultralytics.data.build.InfiniteDataLoader'>
```

Data loaders are *iterables*. That is, you can put them in a `for` loop to load data one batch at a time. We just want to read one batch from it, though.

Task 3.5.5: Load one batch from `loader` into the variable `batch`. You can do this by constructing a `for` loop over `loader` and calling `break` inside the loop, so that it only runs once.

```
In [15]: batch = next(iter(loader))
print(type(batch))
<class 'dict'>
```

A more advanced way to accomplish this same thing is: `batch = next(iter(loader))`

We get back a dictionary. (What a surprise!) Let's explore what's in this structure.

Task 3.5.6: Print out the keys in `batch`.

```
In [19]: print(batch.keys())
dict_keys(['im_file', 'ori_shape', 'resized_shape', 'img', 'cls', 'bboxes', 'batch_idx'])
```

Task 3.5.7: Print out the shape of the `img` value.

```
In [23]: print((batch["img"].shape))
torch.Size([8, 3, 640, 640])
```

The dimension of 3 represents the color channels. The dimension of 640 are the width and height. So what does the dimension of 8 represent?

You can get a clue by reviewing the call to `model.train`. We set a batch size of 8. This tensor thus represents eight training images.

Task 3.5.8: Print out the shape of the `bboxes` value.

```
In [24]: print(batch["bboxes"].shape)
```

That seems like a lot of bounding boxes for one image, so these must be the boxes for all of the images in the batch.

The exact number of bounding boxes will depend on the random batch that got delivered to you. If you re-run the cell that creates `batch`, you'll find that you get another number here.

The image index in the batch that the box corresponds to is given in the `batch_idx` value.

```
In [25]: print(batch["batch_idx"])
```

Thus, we can select the bounding boxes for a particular image in a batch by finding the rows that correspond to a particular batch index value. This is implemented for us in the following function, which will plot the bounding boxes on top of the image.

```
In [26]: def plot_with_bboxes(img, bboxes, cls, batch_idx=None, index=0, **kw):
    """Plot the bounding boxes on an image.

    Input: img      The image, either as a 3-D tensor (one image) or a
           4-D tensor (a stack of images). In the latter case,
           the index argument specifies which image to display.
           bboxes   The bounding boxes, as a N x 4 tensor, in normalized
                    XYWH format.
           cls      The class indices associated with the bounding boxes
                    as a N x 1 tensor.
           batch_idx The index of each bounding box within the stack of
                     images. Ignored if img is 3-D.
           index    The index of the image in the stack to be displayed.
                     Ignored if img is 3-D.
           **kw     All other keyword arguments are accepted and ignored.
                     This allows you to use dictionary unpacking with the
                     values produced by a YOLO DataLoader.

    """
    if img.ndim == 3:
        image = img[None, :]
        index = 0
        batch_idx = torch.zeros((len(cls),))
    elif img.ndim == 4:
        # Get around Black / Flake8 disagreement
```

```
    indp1 = index + 1
    image = img[index:indp1, :]

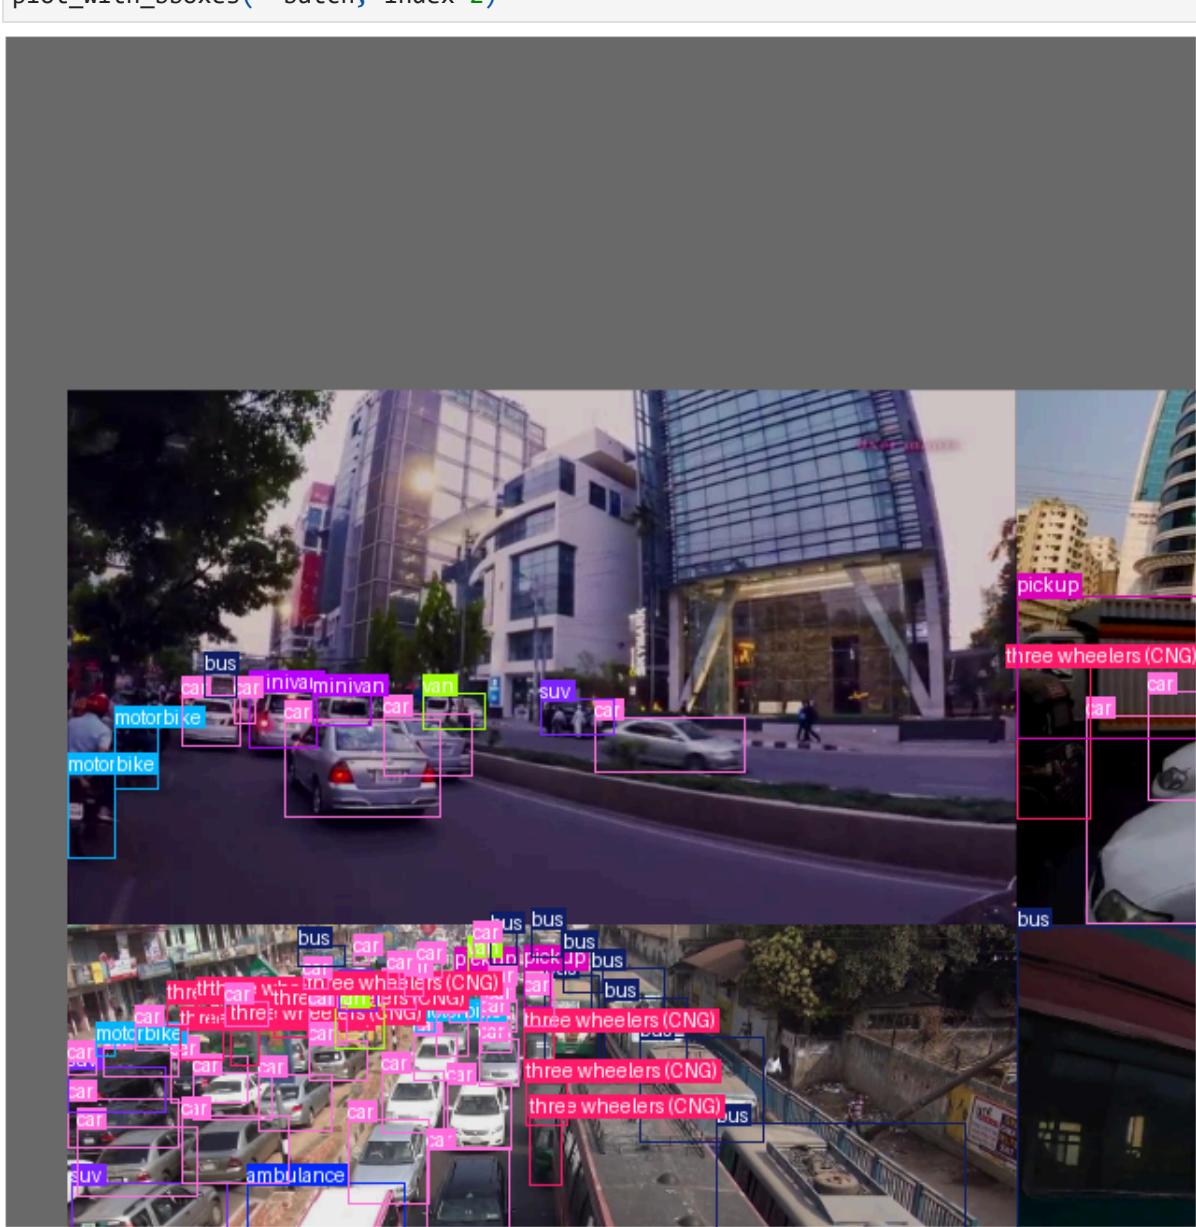
    inds = batch_idx == index
    res = ultralytics.utils.plotting.plot_images(
        images=image,
        batch_idx=batch_idx[inds] - index,
        cls=cls[inds].flatten(),
        bboxes=bboxes[inds],
        names=CLASS_DICT,
        threaded=False,
        save=False,
    )

    return Image.fromarray(res)
```

Task 3.5.9: Plot the image and bounding boxes for index 0 of this batch.

In [33]: `plot_with_bboxes(**batch, index=2)`

Out[33]:



That's ... weird looking. It's not what our original images look like, is it?

If it's not weird looking, try looking at another index. Eventually you'll find something weird looking!

The file names from the batch are stored in the `im_file` key. We can use that to look up the original image associated with this index and see what it looks like.

Task 3.5.10: Display the original image file for this index.

In [30]: `Image.open(batch["im_file"])[2])`

Out[30]:



Comparing the two, we can see that the original image was distorted and combined with other images before being loaded into the YOLO model. The YOLO model applies a number of augmentation steps by default. (You can take a look at [all of the augmentation settings](#) in YOLO.) This increases the diversity of training images, which should help the model generally.

Data Augmentation with Torchvision

If you're training a YOLO model, it's generally best to use the built-in augmentation setting. But in other cases, you may need to implement an augmentation system yourself.

Torchvision makes this easy by providing a number of augmentation transforms in its transforms version 2 (v2) module.

To demonstrate this, we'll load a sample image. The code below will get the file paths for `01.jpg` and its associated label file. (It's written so that it works whether the image ended up in the training or validation split.)

```
In [34]: yolo_base = pathlib.Path("data_yolo")
sample_fn = next((yolo_base / "images").glob("*/01.jpg"))
sample_labels = next((yolo_base / "labels").glob("*/01.txt"))

print(sample_fn)
print(sample_labels)
```

data_yolo/images/val/01.jpg
data_yolo/labels/val/01.txt

Task 3.5.11: Load the image with PIL.

```
In [35]: sample_image = Image.open(sample_fn)

sample_image
```

Out[35]:



Task 3.5.12: Convert the image to a tensor. In the transforms version 2 module, this can be done with the confusingly-named `ToImage` transform.

```
In [36]: sample_torch = v2.ToImage()(sample_image)

print(sample_torch.shape)
```

torch.Size([3, 800, 1200])

The bounding boxes are stored in the label file. Let's take a look at the first five lines to remember what it looks like.

```
In [37]: !head -n5 $sample_labels
```

```
4 0.8 0.74375 0.2116666666666667 0.5125
4 0.7995833333333333 0.424375 0.0975 0.13875
4 0.7995833333333333 0.33 0.0841666666666667 0.0575
13 0.66375 0.595625 0.0591666666666666 0.15875
13 0.66875 0.483125 0.0425 0.05625
```

Each line represents a bounding box. The first element is the class index. This is followed by the x and y coordinates of the box center, the width, and the height.

Task 3.5.13: Load the bounding box data into a variable named `label_data`. It should be a list of the bounding boxes. Each bounding box will itself be a list of five strings in the same order they are in the file. Don't worry about converting the strings to numbers yet.

```
In [39]: # Load the data into `label_data`
with open(sample_labels, "r") as f:
    label_data = [row.split() for row in f]

label_data[:5]
```

```
Out[39]: [['4', '0.8', '0.74375', '0.2116666666666667', '0.5125'],
           ['4', '0.7995833333333333', '0.424375', '0.0975', '0.13875'],
           ['4', '0.7995833333333333', '0.33', '0.0841666666666667', '0.0575'],
           ['13', '0.66375', '0.595625', '0.0591666666666666', '0.15875'],
           ['13', '0.66875', '0.483125', '0.0425', '0.05625']]
```

Task 3.5.14: Create a tensor containing the class indices. For compatibility with our plotting function it should be a $N \times 1$ tensor.

```
In [44]: classes = torch.tensor([[int(row[0])]] for row in label_data])

print("Tensor shape:", classes.shape)
print("First 5 elements:\n", classes[:5])

Tensor shape: torch.Size([30, 1])
First 5 elements:
tensor([[ 4],
       [ 4],
       [ 4],
       [13],
       [13]])
```

Task 3.5.15: Load the bounding box coordinates into a $N \times 4$ tensor.

```
In [62]: bboxes = torch.tensor([[float(el) for el in row[1:]] for row in label_data])

print("Tensor shape:", bboxes.shape)
print("First 5 elements:\n", bboxes[:5])
```

```
Tensor shape: torch.Size([30, 4])
First 5 elements:
tensor([[0.8000, 0.7437, 0.2117, 0.5125],
       [0.7996, 0.4244, 0.0975, 0.1388],
       [0.7996, 0.3300, 0.0842, 0.0575],
       [0.6637, 0.5956, 0.0592, 0.1587],
       [0.6687, 0.4831, 0.0425, 0.0562]])
```

All of these coordinates are normalized by the width or height, as appropriate. This won't work with transformations like rotation, which need the same units used on each axis.

```
In [63]: sample_width, sample_height = sample_image.size

scale_factor = torch.tensor([sample_width, sample_height, sample_width, sample_height])

bboxes_pixels = bboxes * scale_factor

print("Tensor shape:", bboxes_pixels.shape)
print("First 5 elements:\n", bboxes_pixels[:5])
```

```
Tensor shape: torch.Size([30, 4])
First 5 elements:
tensor([[960.0000, 595.0000, 254.0000, 410.0000],
       [959.5000, 339.5000, 117.0000, 111.0000],
       [959.5000, 264.0000, 101.0000, 46.0000],
       [796.5000, 476.5000, 71.0000, 127.0000],
       [802.5000, 386.5000, 51.0000, 45.0000]])
```

Task 3.5.16: Convert the bounding box coordinates to pixel units.

In order for the transformations to know how to transform the bounding boxes, they need to know that the coordinates represent the centers and dimensions of the boxes. This is done by creating a special `BoundingBoxes` tensor. This type has a `format` attribute. By setting this to `"CXYWH"`, we're telling it that the columns represent the Center X coordinate, the Center Y coordinate, the Width, and the Height. This tensor also is given the size of the image, so it doesn't need to look that up for transformations.

```
In [64]: bboxes_tv = torchvision.tv_tensors.BoundingBoxes(
    bboxes_pixels,
    format="CXYWH",
    # Yes, that's right. Despite using width x height everywhere
    # else, here we have to specify the image size as height x
    # width.
    canvas_size=(sample_height, sample_width),
)

print("Tensor type:", type(bboxes_tv))
print("First 5 elements:\n", bboxes_tv[:5])
```

```
Tensor type: <class 'torchvision.tv_tensors._bounding_boxes.BoundingBoxes'>
First 5 elements:
tensor([[960.0000, 595.0000, 254.0000, 410.0000],
       [959.5000, 339.5000, 117.0000, 111.0000],
       [959.5000, 264.0000, 101.0000,  46.0000],
       [796.5000, 476.5000,  71.0000, 127.0000],
       [802.5000, 386.5000,  51.0000,  45.0000]])
```

Let's double check that we did all of those conversions correctly. Do the bounding boxes line up with the correct objects?

In [65]: `plot_with_bboxes(sample_torch, bboxes, classes)`

Out[65]:



If everything looks good, we'll introduce some transformations. The first one will be a horizontal flip. Many everyday objects have bilateral symmetry (or nearly so), so a flipped image will still have the same object classes in it. This makes a horizontal flip a good data augmentation transformation.

(In contrast, up/down symmetry is less common. A vertical flip is generally not as useful, unless you need to recognize upside-down objects.)

The transforms version 2 module has a `RandomHorizontalFlip` transformation. This takes the probability of a flip as an argument.

Task 3.5.17: Use the `RandomHorizontalFlip` transformation to flip the sample image. Set `p=1` to ensure that the flip happens.

```
In [66]: flipped = v2.RandomHorizontalFlip(p=1)(sample_torch)

plot_with_bboxes(flipped, bboxes_tv, classes)
```

Out[66]:



The image has flipped, but the bounding boxes are still in their original locations. Note that the bus is now in the bottom left of the image. Its bounding box is still at the bottom right, and it now contains some asphalt, planters and trees. If we fed this into a model, it would make the model worse, by confusing it as to what a bus looks like.

So, we need to transform the bounding box coordinates consistent with the image transformation. The Torchvision version 2 transformations can take multiple arguments. They perform the same transformation on all of the arguments, returning a transformed version of each. They also understand how to correctly transform the `BoundingBoxes` tensors, depending on their type.

Task 3.5.18: Use `RandomHorizontalFlip` to flip both the sample image and its bounding boxes. Check that they line up correctly now.

```
In [68]: flipped, flipped_bboxes = v2.RandomHorizontalFlip(p=1)(sample_torch, bboxes_tv)

plot_with_bboxes(flipped, flipped_bboxes, classes)
```

Out[68]:

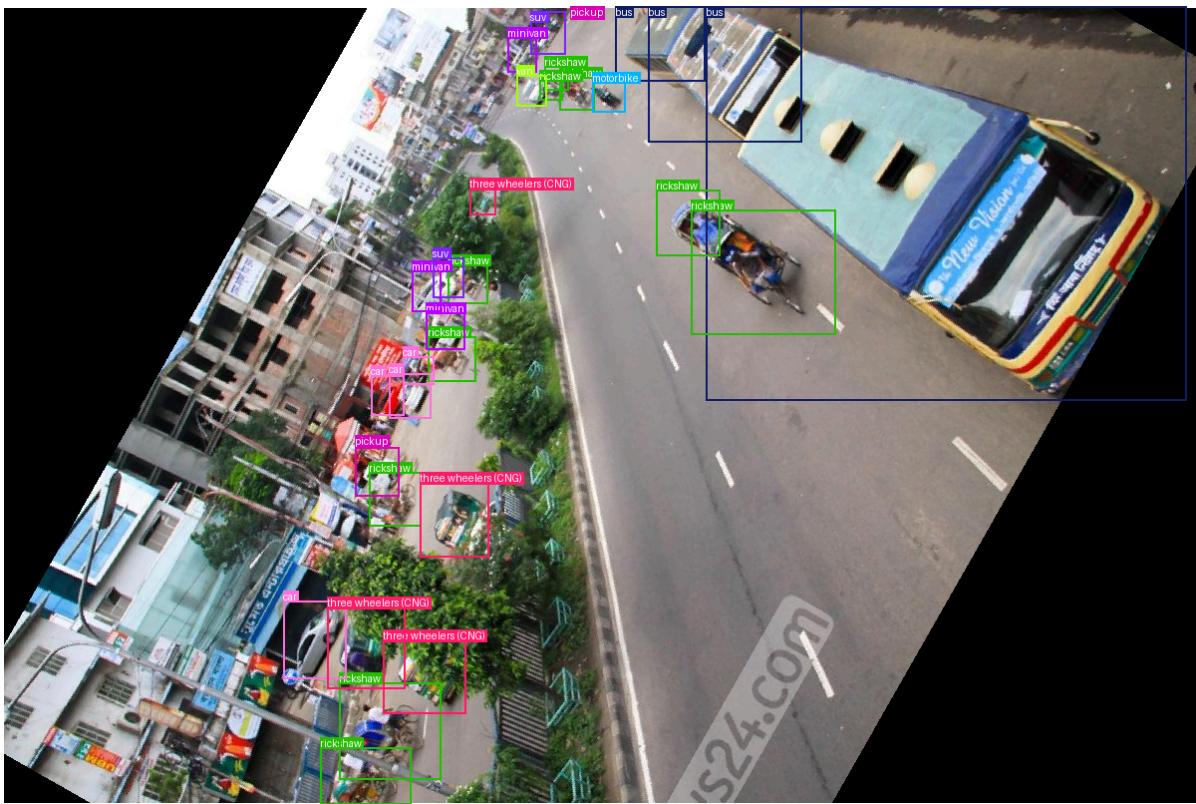


Task 3.5.19: Apply the `RandomRotation` transformation. This takes an argument of the maximum number of degrees to rotate the image. Set it to 90.

```
In [72]: rotated, rotated_bboxes = v2.RandomRotation(degrees=90)(sample_torch, bboxes_tv)

plot_with_bboxes(rotated, rotated_bboxes, classes)
```

Out[72]:



Multiple augmentation techniques can be chained together to produce even more diversity in the training images. Within Torchvision, this can be accomplished by the `Compose` element.

Task 3.5.20: Create an augmentation pipeline that combines the `RandomHorizontalFlip` with the `RandomRotation`. This time, set the probability of the flip to 50%.

```
In [75]: transforms = v2.Compose(
    [
        v2.RandomHorizontalFlip(p=0.5),
        v2.RandomRotation(degrees=90)
    ]
)

transformed, transformed_bboxes = transforms(sample_torch, bboxes_tv)

plot_with_bboxes(transformed, transformed_bboxes, classes)
```

Out[75]:



There are a large number of transformations that can be used for data augmentation. Scroll through [the documentation](#) to get a view of the range of possibilities.

In addition to the transforms we've already used, note:

- `RandomResizedCrop` will randomly crop the image down, and then it resizes the output to a set dimension.
- `ColorJitter` can randomly adjust the brightness, contrast, saturation, and hue of the image, within specified ranges.

Task 3.5.21: Create an augmentation pipeline that applies several of these transformations. Choose reasonable values for the parameters. Check that the bounding boxes are transformed correctly through this.

There's no right answer. A good choice of augmentations depends heavily on the problem you're trying to model.

```
In [84]: transforms = v2.Compose(
    [
        v2.RandomHorizontalFlip(p=0.5),
        v2.RandomRotation(degrees=90),
        v2.RandomResizedCrop(size=(400, 500)),
        v2.ColorJitter()
    ]
)

transformed, transformed_bboxes = transforms(sample_torch, bboxes_tv)
plot_with_bboxes(transformed, transformed_bboxes, classes)
```

Out[84]:



You can run the transformation several times to see the different types of images that result. This greater diversity of training images will help models learn to generalize instead of memorizing during training.